

Aadarsh Dubey

8563906369 | addydubey321@gmail.com | [linkedin.com/in/aadarsh-dubey](https://www.linkedin.com/in/aadarsh-dubey) | github.com/AadarshDubey

PROFESSIONAL SUMMARY

AI Engineer with 2 years of experience building Retrieval-Augmented Generation (RAG) pipelines, multi-agent and agentic AI systems, and full-stack AI applications for financial and enterprise use cases. Experienced with LangGraph, LangChain, Azure AI Search, FastAPI, and Next.js, with hands-on work across vector search, LLM evaluation, microservices, Docker, and CI/CD.

EDUCATION

Vellore Institute of Technology, Bhopal Aug. 2021 – Jul. 2025
Bachelor of Engineering in Computer Science and Engineering CGPA: 8.56/10

SKILLS

Languages: JavaScript (ES6+), TypeScript, Python, SQL
Frontend: React.js, Next.js
Backend: FastAPI, Flask, Node.js, Express.js, PostgreSQL, Redis, MongoDB, REST API, Microservices
AI & ML: LangChain, LangGraph, CrewAI, RAG, Agentic AI, LLM Fine-tuning, LLM Evaluation, Prompt Engineering
Vector Databases: ChromaDB, FAISS, Pinecone, Azure AI Search, Vector Search
DevOps & Tools: Docker, CI/CD (GitHub Actions), Git, VS Code

WORK EXPERIENCE

GenAI Engineer Jun. 2025 – Present
LTM

- Architected a LangGraph-based RAG system on BlueVerse using Azure AI Search, replacing manual data lookups and reducing retrieval time by ~60% based on internal benchmarking.
- Engineered a Crawl4AI ingestion pipeline to automatically collect 100+ financial records per run from Yahoo Finance, transforming a previously manual data collection workflow into an automated RAG knowledge-base pipeline.
- Developed an entity-matching pipeline for an Accounts Payable automation platform using Agno, PaddleOCR, and OLMOCR to perform 3-way and 4-way matching across invoices, POs, GRNs, and contract terms; modeled ~1,100 Neo4j nodes and ~3,700 relationships to enable multi-hop matching and discrepancy detection.
- Led standardization of Agno- and LangGraph-based service scaffolding across BlueVerse projects, introducing structured logging and Arize Phoenix for LLM tracing, observability, and debugging.

LLM Evaluation Engineer Dec. 2024 – Jun. 2025
Outlier — Freelance

- Evaluated 10+ LLMs across 100+ domain-specific tasks using structured evaluation criteria, identifying recurring reasoning, factuality, and instruction-following failures to support downstream model improvement.
- Reviewed 200+ AI-generated code solutions for correctness, security, and performance, identifying bugs and bottlenecks in ~1/3 of submissions and providing actionable human feedback for model improvement.
- Designed 400+ adversarial prompts targeting hallucinations, reasoning failures, and edge cases, surfacing errors in ~20% of evaluations while performing domain-specific validation across cryptocurrency and healthcare tasks.

PROJECTS

CRX-PIPELINE | Crypto Market Intelligence System Feb. 2026

- Designed a RAG pipeline ingesting crypto market data into ChromaDB, indexing 1K+ records for semantic search and trend analysis.
- Used Redis caching and a message queue(ARQ) to cut repeated-query latency to under 200ms under 10 concurrent simulated requests.
- Containerized services with Docker and configured CI/CD via GitHub Actions, reducing deployment steps from manual multi-stage to a single automated workflow.
- Shipped a FastAPI and Next.js dashboard with search, summaries, and pricing insights for traders.

SoW Assistant | AI Contract Auditing System Dec. 2025

- Implemented a three-agent pipeline (Parser, Risk, Quality) that flagged ~85% of intentionally-vague clauses across 10+ test SoW documents.

- Automated contract parsing workflows, cutting manual review time from ~3 hours to under 15 minutes per fixed-price SoW.
- Generated redline suggestions and structured audit reports across 50+ contract sections to support faster reviewer sign-off.

Recapify | Automated Meeting Analysis

Sep. 2025

- Launched a Flask and LangGraph pipeline converting 60+ minute meeting transcripts into structured summary reports in under 5 minutes.
- Orchestrated multi-step LLM workflows to extract action items, owners, and decisions across 7+ categories with ~90% consistent formatting across test transcripts.
- Integrated Gemini for cloud use and Ollama for local use, avoiding reliance on a single LLM provider.

CERTIFICATIONS

Azure AI Engineer Associate (AI-102): Designing and implementing Azure AI solutions including Azure OpenAI, cognitive services, and AI-powered search

Azure AI Fundamentals (AI-900): Foundational knowledge of AI and machine learning concepts using Microsoft Azure AI services

AWS Cloud Practitioner: Architecting scalable cloud solutions and optimizing cloud deployment and management

Google Cloud Digital Leader: Foundational cloud technologies knowledge for driving digital transformation using Google Cloud services